

Folios — Product Overview

FOLIOS · FOLIOSHQ.COM

LAST UPDATED: 2026-06-19 (F6 – PIP SEMANTIC RECALL; MONDAY 1:1 PACK; CHAT AGENT LOOP)


This is the substantive product read after the [one-pager](#). Covers what Folios does, how it's structured, and what makes the Pip layer different from "AI bolted onto a CRM."

Product framing

Folios is the year-round relationship-management workspace for external account managers. It captures the work of relationship management — meetings, follow-ups, commitments, context — and gives that work an intelligence layer that gets better over time.

The product is named **Folios** (plural). The GitHub repo is **Folio** (singular) for historical reasons.

The three modules under the Folios umbrella



- **Folios** — the main app. Accounts, contacts, cadences, meetings, open items, Pip AI. The everyday surface.
- **Lanyard** — conference-specific module. Pulls into focus during trade shows (schedule, partner profiles, team chat, in-person notes). Feeds notes back to Folios after the conference.
- **Gauge** — project & task management module. Tracks commitments from meetings into trackable projects (discrete and standing). Will expand to a full PM tool used by cross-functional product/ops teams.

All three share the same Supabase backend, the same auth, the same Pip.

Core capabilities

Workspaces

Folios organizes accounts into three workspace types, all sharing the same underlying model:

- **Accounts** — customer-facing relationships (standard, MSO, shop)
- **Departments** — internal teams (sales, ops, product, etc.)
- **Partners** — vendor / supplier / channel relationships

The conditional UI changes per type (Accounts show revenue / tier; Partners show agreement terms; Departments hide both). Pip's context also branches per type so it asks the right questions about each.

Sub-accounts are supported via `parent_account_id` (e.g., parent MSO holds child shops; parent customer holds child divisions).

Account detail

Each account gets its own command center:

- **Overview** — health (auto-computed), context notes (read by Pip), follow-up due, recent meetings, recent updates, systems/tools they use (chips, populated via approved Pip suggestions)
- **Cadences** — recurring meeting schedules
- **Meetings** — all conversations with the account
- **Items** — open action items
- **Contacts** — people on the account
- **Projects** (Gauge) — active deliverables
- **Updates** — change events (catalog/pricing/integration/etc.) for revenue-impact analysis
- **Shops** (MSO type only) — child shops

Cadences

A cadence is a recurring meeting schedule attached to an account. Define the frequency (weekly / biweekly / monthly / quarterly), day of week, and optional time. Folios:

- Auto-computes the next due date
- Sends in-app + browser-notification reminders 30m / 5m / start
- Auto-advances when a conversation is logged tied to it
- Surfaces the Cadence Hub when due

Cadences can also be **account-less** — a recurring 1:1 with your manager or a teammate, or an internal/leadership meeting (`cadence_scope='person'`). Action items from those meetings persist as **leadership tasks** — your own to-dos, filed under the 1:1 rather than forced onto a customer account — and show in that cadence's hub.

Cadence Hub

Pre-call command center for any account/cadence. Opens to:

- **Pip brief** — short pre-call read on the relationship state

- **Big Start Meeting button** — auto-creates a draft and opens full-screen meeting mode
- **Active Gauge projects** — expandable inline
- **Open items** — what's hanging
- **Follow-ups** — what's promised
- **History** — every prior conversation

Mobile uses a 4-tab segmented control (Notes / History / Tasks / Follow-ups) for compactness.

Full-screen meeting mode

Distraction-free conversation capture. Opens over the global chrome:

- Top bar with End & Summarize
- **Split-screen layout** (desktop): projects panel on the left — each project card carries its **own meeting-note field** once marked discussed — general notes on the right. Open items and people sit in collapsible sections under the projects, including inline add-contact so new people met mid-meeting are captured on the spot.
- **Mobile:** a Notes / Projects toggle gives each side of the split the full screen, one at a time.
- Per-project notes are stored separately (`project_notes`) so Pip knows exactly which notes belong to which project and routes that project's action items to it precisely.
- Viewport-filling notepad with bullet-paste preservation
- Autosaves every 1.5s (general and per-project notes alike)
- ESC to close; un-discussing a project with typed notes asks before discarding them

On End & Summarize, Pip generates a structured plan (action items, project updates, follow-ups) and the user reviews it in the [summarize-preview modal](#) before applying.

Discussed signal: during a meeting the user can tap any project card or open-item row in the sidebar to flag it as "discussed this meeting." Flagged cards show a teal ♦ Discussed chip. As the user types notes, the sidebar also auto-highlights any project or

item whose title appears in the notes with a subtle glow — giving visual confirmation Pip will connect the dots before hitting summarize. Flagged IDs are passed to Pip at summarize time so Pip strongly prefers updating or closing those items instead of creating new duplicates. Rows in the plan modal that touch a flagged project or item show a ♦ Discussed badge so the user can see Pip acted on the signal.

Scheduling future one-off meetings

In addition to recurring cadences, you can pin a single upcoming meeting to a specific date/time. From the Calendar view:

- Click any empty day to open the schedule modal, or use the **+ Schedule Meeting** button in the header
- Pick the account (required), date, time, method, and an optional agenda note
- The meeting appears on the Calendar, Week, and List views with a distinct ♦ chip (distinguished from cadence events)
- Reminders fire at 30m / 5m / start just like cadence meetings
- Today's scheduled meetings surface on the Home screen in a "Scheduled Today" card
- Opening the card on or after meeting day flips it into a live draft and opens the full-screen meeting mode so the conversation can be captured and summarized normally

Logging conversations (non-cadence)

For ad-hoc calls outside a cadence, the global "+ Conversation" button opens a modal: pick the account, pick the method, pick the date, then open the same full-screen meeting mode. The conversation files under the account's Meetings tab tagged **AD-HOC** .

Open items

Lightweight action-item tracking per account. Items can be:

- Manually added

- Auto-created by Pip from meeting notes (with `pip_created_at` marker so the V2 brain can detect post-creation edits)
- Carried forward across cadence meetings (open items show up in the next meeting's prep view)

Items support due dates, assignees (org members), and completion.

Quick Tasks

A floating tray for fast-capture tasks that don't belong to a specific account. 2-second log without leaving whatever screen you're on. Same underlying row shape as project tasks.

Update Calendar

Manual change-event log per account. Categories: catalog, pricing, integration, product launch, training, promo, external event, other. Each entry has a date, title, description, owner, observed impact.

The value: when revenue dips on an account, you can correlate against the update timeline to see what changed and when. (Revenue surfaces themselves have been removed from the personal-mode build — see "Ripped" in CLAUDE.md — but the change-event log remains valuable.)

Inactive & Merge

No hard deletes for accounts or users. Everything is reversible:

- **Inactive flag** — soft archive. Greyed out in lists, excluded from Pip's "needs attention" counts. Can be reactivated.
- **Merge into...** — for acquisitions. Re-parents every child row (meetings, items, contacts, cadences, projects) from source → target, then marks source inactive. Linked back so history is preserved.

Pip — the AI field analyst

Pip is the differentiator. Embedded in every meeting, every account view, every brief.

Pip's visual form

Pip has a distinct physical presence across every screen: two glowing hex-lattice spheres — a head and a smaller tail — suspended inside a slowly turning hexagonal ring. The whole figure breathes together on a 2.4-second cycle: the ring swells, the sphere hexes open and close, and the core glow pulses, all in sync. The 3D scene renders via an SVG animation loop and re-skins automatically when the accent token changes (Work mode teal → Life mode dusty orange). Smaller instances ($\leq 32\text{px}$) stay as the classic two-circle form since the hexes mush at small sizes. The design is locked and enforced by an automated CI test — parameters cannot drift without a visible failure.

Pip's personality

Loyal, slightly anxious field analyst. Same voice across all surfaces — meeting summaries, briefs, chat answers, error recovery messages. Personality is intentional: feels like a real teammate, not a generic chatbot.

Pip's modes

Mode	What it does	When it fires
summarize	Turns meeting notes into a structured plan (new items, updated items, closed items, project updates, follow-ups)	End & Summarize on a meeting
brief	Pre-call brief for an account (relationship state, what's at stake)	"Brief Me" button on Cadence Hub
ask	Conversational chat — ask anything about your accounts	Pip view (chat surface)
plan / extract	One-shot extraction of action items from short notes	Quick Touchpoint modal

Model tiering. Pip runs the right model for each task. The reasoning-heavy, user-facing, low-frequency surfaces — Ask Pip chat, meeting summarize, portfolio question generation, the daily brief, profile synthesis, and the QBR — run on **Sonnet 4.6**. High-volume or mechanical surfaces — per-account Brief Me, follow-up email drafting, terminology extraction, memory compression, and the leadership readout — run on **Haiku 4.5**. Each Sonnet surface has an environment override so its tier can be re-dialed without a code change. (See [ai-governance.md](#) for the full matrix.)

How Pip writes. Every Pip text surface — chat, the daily brief, meeting summaries, Brief Me, the cadence pre-call brief, and the QBR — renders structured output: a short headline, labeled sections, and scannable bullets with the important names in bold, instead of a wall of prose. Priority sections can carry one of Pip's own inline status glyphs (a small on-brand icon set — needs-now, keep-an-eye, good-news, cross-account-pattern, done) in place of generic emoji, so a brief reads at a glance. Email drafts stay plain prose (no markup) since they're meant to be sent as-is.

Autonomous Operator — Pip works the book on demand

Pip is not only pull-triggered. A manual operator pass sweeps the portfolio on demand, so the user can run it when they want a fresh read rather than staring at a dashboard to interpret.

How to run: tap "Run Pip's pass" on the Home screen. The pass takes 20–40 seconds and updates the operator report when complete. The Vercel nightly cron was retired June 2026 (it ran every morning regardless of whether the app was opened, burning tokens on quiet days).

- **Signal-gated sweep.** The operator job reviews the book and runs a deep per-account pass *only on the accounts that moved* since the last run (new activity, or a health/at-risk signal). Accounts that didn't change are skipped — cost and effort scale with what actually changed, not with portfolio size. The per-account passes are capped per run.
- **Materialized operator state.** For each worked account Pip writes a durable state object: a read of the situation, the risks, a **pre-drafted follow-up email** where one is warranted, proposed task/project moves, a pre-built cadence agenda, and a "what changed since last run" delta. This is stored, not regenerated on every screen open — surfaces read it instead of making their own model call.

- **The operator report.** Those per-account reads roll up into one prioritized report on the Home screen: the plan for the day, with the drafted follow-ups ready to review and send. It replaces the live daily brief on days the pass ran ("here's what's happening" becomes "here's what's happening, and I drafted the first pass of it").
- **Surfaced across the app, one brain.** The same materialized state is read directly by every relevant surface — no extra model call when a screen opens: the **account screen** shows a full operator panel (situation, "since last run" delta, risks, the drafted email, and proposed moves you approve or dismiss one tap at a time); the **Cadence Hub** shows the pre-built agenda before a call; the **Gauge Pip card** shows proposed moves as a cross-account decision queue. Approving a proposed move creates the task; the proposal is then marked handled so it doesn't reappear.
- **Propose, don't act.** Everything the pass produces is a draft the user approves — nothing is created live and nothing is sent automatically. This keeps the data boundary clean: Pip never reaches outside Folios.
- **Cost scales with change, not size.** The deep per-account pass only runs on accounts that moved since the last run, capped per pass, plus one portfolio roll-up — so each run's spend tracks what actually changed across the book.

The V2 brain — Pip's learning loop

This is the differentiator that nothing else does. Pip doesn't just generate text; Pip **LEARNS** from how the user reacts to its output.

Four learning surfaces:

0. **User profile (folio_user_profile + profile_prose).** At first login, Pip interviews the user with 5 questions (role, company, portfolio, goals, communication style). Answers are synthesized into a 4-8 sentence **profile_prose** narrative injected into every Pip response — chat and summarize equally. Pip now knows who it's talking to, not just what accounts they have. Soft-gated: interview is skippable and resumable. Existing users see a dismissible HomeView card. Cost: ~\$0.002 once.
Phase 2 — drip questions. After onboarding, Pip keeps learning via a "Pip's Curious" card on the Home screen — one question at a time, inline textarea,

never a modal. Soft cap of 5 per rolling 24h with no skip cooldown, so you can power through several in a sitting. Every question comes from real observed data — there is no generic "get-to-know-you" filler (that was removed; the well staying silent beats the well being fake). Sources:

- **Structural gaps (zero LLM cost):** contacts who appear in ≥ 3 meetings but have no role recorded, active accounts past 30 days with no objective, and empty profile slots post-onboarding.
- **Terminology lane (Lane C):** a daily Haiku scan of recent meeting notes surfaces proper nouns / brand names / codenames that appear ≥ 3 times but aren't a known account, contact, or glossary term. Each question is account-anchored ("you keep mentioning Fuse5 around John's Auto Parts — what is it?"). Answers write to `folio_pip_facts` so future briefs know the company's vocabulary automatically.
- **Portfolio generator (Lane D):** a low-frequency Sonnet pass reasons across the whole portfolio and writes a few genuinely insightful questions in Pip's voice. It self-skips (a DB count, no model call) whenever the queue already holds ≥ 5 , so it only spends when the queue has actually drained.

Answering ≥ 3 drip questions since the last synthesis triggers a background re-synthesis so `profile_prose` stays current. Settings → "Pip's Questions" has a global pause toggle and a completeness meter. Cost: gap detection = \$0; terminology + generator $\approx$ \$0.01-0.02/month; re-synthesis $\approx$ \$0.004 per batch.

1. **Correction log (`pip_correction_log`).** Every time the user declines a proposed action, edits the text of a Pip-created item, reassigns a task, or rejects a routing decision, the correction is captured. Pip reads the last 10 corrections per account on every summarize call and is explicitly told "Chris has corrected these before — don't repeat."
2. **Lessons learned (`pip_account_state.lessons_learned`).** Every ~ 5 meetings, a cheap compression pass distills the recent correction log into a stable paragraph stored per account. Pip reads this paragraph instead of the raw log, keeping the read-back context small but the institutional memory permanent.
3. **Assignment hints (`pip_assignment_hints`).** When the user reassigns a task in the summarize-preview modal (e.g., changes assignee from "Mark" to "Sara" for invoice work), the pattern is captured. Pip routes the next similar task

to Sara automatically. Hints scope: account-specific first, then org-wide ("Sara does invoice work everywhere").

4. **Glossary (`pip_glossary`).** Users can teach Pip terms, acronyms, and aliases (e.g., "ProParts was the legacy name for KSI Collision"). Pip respects glossary aliases when matching account names in notes.
5. **Structured suggestions ("Pip proposes, you approve").** When a drip question is *about* something structured — a contact's role, an account's objective, or a tool/system the account uses — the intent is attached to the question when it's created (Pip already knows which account or contact it's asking about). On answer, the card shows a pre-checked "Also save to..." toggle: keep it checked and the answer is written to the real field (`folio_contacts.title` , `folio_accounts.objective` , or the account's `systems` list); untick it to keep the answer as a plain fact. One tap, human in the loop — nothing mutates silently, and it never touches health or tier. The "systems" an account uses (e.g. "Fuse5 is their IMS") surface as chips on the account Overview and ride into every per-account Pip surface, so when "Fuse5" later shows up in raw notes Pip knows what it is instead of asking again. No per-answer model call — the structured write is free.

Cross-account routing. Pip can route a task from one account's meeting to a different account's plate. Example: "ACME meeting, but I need to follow up with KSI on the invoice feed" → task created on KSI, source backref points to the ACME meeting. The V2 brain handles the routing decision with a `target_account_id` field on each plan row. User can override in the preview modal.

Internal-meeting prior. When the meeting is on an internal-team account (`account_type = 'internal_team'`), Pip's default flips — it expects tasks to fan out to customer accounts, not stay on the internal account.

Pip — semantic recall (pgvector)

Recency context only ever shows Pip the *latest* N meetings on an account. Some of the most useful context is older — a decision made six months ago, a constraint a contact mentioned once. Semantic recall closes that gap: Pip can pull the most relevant past notes by **meaning**, not date.

How it works. Folios keeps a vector index (`folio_embeddings` , Postgres pgvector) of the user's own content — meeting notes, Pip's meeting summaries, per-project meeting notes, and account updates. A daily background sweep embeds anything new or changed (each source carries a content fingerprint, so nothing is ever re-embedded unchanged — the running cost is effectively zero). When the user asks Pip a question, the question is embedded and the closest past notes are retrieved — account-scoped when the question is about one account, or across the whole book for "what did we ever decide about X" questions — and folded into Pip's context through the same shared context builder every Pip surface reads from, so recall reaches chat by construction.

Privacy & data line. Recall is strictly the user's own content, RLS-scoped to them and account-scoped by default (the retrieval function runs under the caller's identity with an explicit owner check). Only data-line-clean text is embedded — no revenue, volumes, or rosters. Recall stays inert until an embeddings key is configured; without it, Pip simply falls back to recency.

Pip — the summarize-preview modal

After Pip generates a meeting plan, the user reviews it before anything is applied. The modal shows every proposed action as a row with:

- Checkbox (default checked — uncheck to decline)
- Editable title (fix wording before applying)
- Account chip (override the routed account if Pip got it wrong)
- Assignee dropdown
- Due date picker
- "See source" link → expands the slice of the user's notes that triggered the row (editable; edited excerpts feed back as correction reasons)
- "Add missed item" button at the bottom for things Pip didn't catch

Grouped into Changes / New / Skipped. Yellow dot on low-confidence rows. Apply runs the selected rows through real DB writes. Cancel preserves the summarized meeting but applies nothing.

Pip — cost optimization

- **Model selection:** claude-haiku-4-5-20251001 for default Pip operations. Sonnet only for high-synthesis paths (Brief Me).
- **Stacked prompt caching:** 4 cache breakpoints (system / glossary / roster / items+tasks). Multi-call sessions get ~70% cache hit rate.
- **Trivial-draft short-circuit:** drafts under 100 characters never call the API.
- **Per-user usage telemetry:** `folio_pip_usage` tracks every call (tokens, mode, timestamp); visible to user in Settings.
- **Rate limiting:** 20 req/min per user.

Pip — voice

- Speech-to-text via browser Web Speech API (free, browser-native).
- Text-to-speech via browser SpeechSynthesis today (free, but the voice quality is poor).
- Premium voice (OpenAI / Cartesia / ElevenLabs) on the wishlist.

Pip — portfolio intelligence ("chief of staff")

Pip sees your entire portfolio at once, not just one account at a time.

- **Daily account state snapshots** (`folio_account_snapshots`): computed once per day, client-side, from real signals — health status, days since last contact, open and overdue item counts, active and stuck Gauge project counts. Zero LLM cost; pure data derivation.
- **Daily brief card on the Home screen:** one call per day, cached for the calendar day. Pip synthesises a morning read across the whole portfolio — what's due, what's at risk, recent wins — factoring in workload (overdue tasks, commitments, today's cadences), cross-account themes, tone trend, health momentum, and **off-cadence accounts vs. their own rhythm** (e.g. "you usually meet every ~3 weeks; it's been 45 days" — a personal baseline, not a generic cold threshold).
- **Draft-ahead follow-ups:** Home surfaces meetings you summarised a couple days ago that still have no follow-up logged, with the follow-up email Pip

already drafted ready to copy or send — no extra cost.

- **Boss-ready leadership readout** and **pre-call standing agendas** (via the cadence brief) round out the proactive layer.
- **Portfolio state utility (`buildPortfolioState`)**: a compact text block summarising at-risk / watching accounts and stuck projects, injected into cross-portfolio Pip surfaces.

Pip — the Monday 1:1 pack

An auto-assembled prep sheet for your weekly 1:1, so nothing surprises you in front of your boss. It surfaces on the Home screen on Monday (and as a Sunday- evening heads-up) and opens into the full pack inside that 1:1's cadence hub. The pack reads top to bottom:

1. **Pip's read** — one to three sentences framing the week.
2. **Your word** — promised-vs-done: the week's commitments, each marked **Kept / Slipped / Open**, with the account.
3. **Boss's open asks, pre-answered** — Pip pulls the asks from your last 1:1's notes (and the open leadership tasks tagged to that 1:1) and attaches the current status to each, so you walk in with the answer ready.
4. **What moved, by account** — the week's per-account delta: meetings, Gauge status pulses, deliveries.
5. **Who has the ball** — waiting-ons: what you owe vs. what's owed you.

Almost all of it is **deterministic assembly over data you already captured** — zero AI cost and always fresh. Only the read + the boss-ask extraction use a model, folded into **one Sonnet call per week**: the output is cached on the cadence row and regenerated only when the week's content actually changes (the same event-driven gating used for Pip's per-account state), so a quiet week never re-bills. The boss-ask extraction reads what you already type — there is **no extra tagging step**. Data Line Rule applies throughout: asks and statuses are directional, never numbers, and your raw 1:1 notes are read, never rewritten.

Gauge — projects & tasks

Gauge is the project-management module that lives under the Folios umbrella. Phase 1 is AM-facing; Phase 2 will expand to a full cross-functional PM tool.

Project types

Discrete projects. AM-owned, multi-step, has a finish line. The AM is the conductor; individual steps have assignees who may or may not be the AM. Example: an audit with 7 steps, 3 different people involved.

Standing projects. Admin-owned reactive queues. No sequence, no finish line. AMs (and Pip) drop tasks in; admin clears them. Example: an "Invoice Updates" inbox.

A project can be linked to multiple accounts — the account field is a searchable picker (search → pick → chip → search again), so a project that touches several accounts is filed under all of them.

Notes vs. status updates

Each project carries two distinct text surfaces:

- **Notes** — a durable scratchpad for context, decisions, and background that stays put.
- **Status updates** — an append-only, timestamped pulse log. Posting an update prepends a new entry (body + timestamp + author); the expanded card shows the latest pulse ("Updated 2h ago"), and the project edit screen shows the full history. Pip reads the latest update (plus the prior two) so briefs can say "All Star — latest: 'waiting on legal sign-off' (Jun 3)."

Custom fields

Each project defines its own custom field schema (text, longtext, number, date, dropdown, person, checkbox, URL). A "bones" preset provides defaults (Priority, Owner, Submission Date, Due Date, Description, Related Link) when starting blank.

Standing board view

Standing projects render as a kanban board inside the project row. One column per `task_status_columns` entry. Tasks drag between columns; admin processes left-to-right.

My Queue view

Flat task list across all projects the current user is assigned to. Live / Planning / All sub-filters. Optional group-by-project toggle.

Gauge — three lenses (queued, not yet built)

The next major Gauge build introduces three role-based default views set at invite time:

- **AM lens** — your accounts, your projects, your queue
- **Leader lens** — team rollup, who's stuck, what's overdue
- **Admin lens** — flat task queue, due dates, account chip per row

Pip's prompt branches per lens for role-appropriate framing. Same UI shape; different defaults and Pip context.

Full spec lives in CLAUDE.md → section 14 → Gauge V3.

Team & org layer

Multi-tenant ready from day 1:

- `folio_orgs` — org name, plan
- `folio_org_members` — user → org → role
- Roles: `owner | admin | member`
- All user-data tables carry `org_id` (or are linked via account)

- RLS policies scope to org membership where appropriate

Today: single-user pilot, no real org sharing in production. Architecture is ready when team mode lights up.

Work / Life mode

Folios has two top-level modes — **Work** (green palette, account management) and **Life** (dusty-blue palette, personal assistant). A toggle in the desktop rail (bottom-left) and the mobile header switches modes instantly; the app recolors via CSS custom property swap, same mechanism as the light/dark theme.

Work mode is the full account management OS — accounts, cadences, Gauge, calendar, meetings. Unchanged by the Life build.

Life mode (Phase 1, June 2026) surfaces:

- **Upcoming** — appointments and events, with a VIP heads-up ladder (anniversary / spouse's birthday / Christmas get multi-stage escalating nudges: 3wk soft → 1wk → 3d → 1d → day-of; annual recurrence set once)
- **Honey-do list** — open home tasks prioritized by aging + complexity
- **Soccer card** — Man United / Brazil / USMNT news + scores (moved from work Home into Life Home)

Phase 2 (planned) = honey-do brain: Pip as a home-project partner with how-to coaching, materials lists, and step-by-step guidance.

Pip's orb recolors to dusty orange in Life mode (reads **- -accent**). Both modes share one Supabase project and one Pip brain.

Mobile Home — structured hub

The mobile Home screen is an `OperatorHub` layout (since June 2026): a compact Pip glance card (headline chip + count chips + collapsed full read) followed by section cards (Today / This Week / Good News / Pattern) with tinted header strips, count badges, and clean stacked rows. An "On the Calendar" card surfaces same-day meetings. Redundant narrative panels are suppressed when an operator report is active.

Teach Pip — on-demand knowledge-building

Beyond the passive weekly drip question, the user can open a "Catch up with Pip" session anytime (Home card when no question is queued, or Settings → Pip's Questions). The session works through the question queue and offers a "Pip, ask me more →" button that fires a fresh manual generation pass (bypassing the 6-hour throttle) so the session never dead-ends. Deterministic account-anchored fallback questions guarantee at least one question per session even when the model abstains. Every answer folds into `profile_prose` and `folio_pip_facts` within minutes (re-synthesis throttled at 5 min).

Light + dark themes

Full design system supporting both palettes. Tokens defined as CSS custom properties on `<html data-theme="...">`. Pre-mount inline script prevents flash-of-wrong-theme. Settings → Appearance for toggle.

Dark is the canonical palette for layout decisions; light is the spec'd translation. Light-mode-specific behaviors (hover lifts, mark pulses) are scoped via `[data-theme="light"]` so dark renders cleanly.

Mobile / PWA

Folios is mobile-first:

- Installable as a PWA on iOS and Android
 - Offline-capable shell (service worker)
 - Two redundant auto-update paths (controllerchange + version polling)
 - Safe-area-aware layouts
 - 16px minimum font on all inputs (prevents iOS auto-zoom)
 - Bottom nav: Home / Accounts / Calendar / Gauge
 - Settings/diagnostics/etc. tucked into the user menu
-
-

Observability

Self-monitoring built in:

- **folio_errors** — every React render error, network failure, Pip error, unhandled rejection lands here with stack + context. Visible in Settings → Diagnostics.
 - **Auto-recovery** — stale-chunk Lazy import failures (post-deploy edge case) detected and reloaded silently; marked auto-resolved.
 - **Connection status** — UI indicator when Supabase realtime drops.
 - **Render-time `timed()` helper** wraps hot paths for performance diagnostics.
-
-

What's NOT in Folios (deliberate)

- **Pipeline / revenue tracking** — ripped from the personal-mode build. Compliance blocks real revenue ingestion. Schema preserved for future re-build

if corporate-data integration becomes possible. See CLAUDE.md → "Ripped (deliberate simplification)" for the inventory of what was pulled.

- **Email integration** — no inbox sync, no auto-log from email threads. Pip operates on notes the user explicitly captures.
 - **CRM sync** — no Salesforce / HubSpot bidirectional sync today.
 - **Automatic transcription** — meeting notes are user-typed (or voice-dictated). No Otter-style auto-transcribe from a Zoom feed.
-
-

Roadmap snapshot

For the current shipping queue and wishlist, see CLAUDE.md → "Pending Updates" and "Feature Wishlist / Roadmap" sections (internal working doc). A presentation-facing roadmap.md is planned.

Highlights of what's next:

- **Gauge V3** (three role-based lenses + unified `folio_tasks` table
 - project templates + Leader project view) — spec'd, queued.
 - **Premium TTS for Pip** (OpenAI / Cartesia / ElevenLabs) — wishlist.
 - **Pip's weekly clarifying questions** — spec'd, depends on more correction-log data first.
 - **Active session management UI** — for multi-device sign-out.
-
-

Contact

chris.vasconcellos97@gmail.com