

Folios — Product Overview

FOLIOS · FOLIOSHQ.COM

LAST UPDATED: 2026-06-06 (LEADERSHIP TASKS + PIP CHIEF-OF-STAFF FINISH)


This is the substantive product read after the [one-pager](#). Covers what Folios does, how it's structured, and what makes the Pip layer different from "AI bolted onto a CRM."

Product framing

Folios is the year-round relationship-management workspace for external account managers. It captures the work of relationship management — meetings, follow-ups, commitments, context — and gives that work an intelligence layer that gets better over time.

The product is named **Folios** (plural). The GitHub repo is **Folio** (singular) for historical reasons.

The three modules under the Folios umbrella



- **Folios** — the main app. Accounts, contacts, cadences, meetings, open items, Pip AI. The everyday surface.
- **Lanyard** — conference-specific module. Pulls into focus during trade shows (schedule, partner profiles, team chat, in-person notes). Feeds notes back to Folios after the conference.
- **Gauge** — project & task management module. Tracks commitments from meetings into trackable projects (discrete and standing). Will expand to a full PM tool used by cross-functional product/ops teams.

All three share the same Supabase backend, the same auth, the same Pip.

Core capabilities

Workspaces

Folios organizes accounts into three workspace types, all sharing the same underlying model:

- **Accounts** — customer-facing relationships (standard, MSO, shop)
- **Departments** — internal teams (sales, ops, product, etc.)
- **Partners** — vendor / supplier / channel relationships

The conditional UI changes per type (Accounts show revenue / tier; Partners show agreement terms; Departments hide both). Pip's context also branches per type so it asks the right questions about each.

Sub-accounts are supported via `parent_account_id` (e.g., parent MSO holds child shops; parent customer holds child divisions).

Account detail

Each account gets its own command center:

- **Overview** — health (auto-computed), context notes (read by Pip), follow-up due, recent meetings, recent updates, systems/tools they use (chips, populated via approved Pip suggestions)
- **Cadences** — recurring meeting schedules
- **Meetings** — all conversations with the account
- **Items** — open action items
- **Contacts** — people on the account
- **Projects** (Gauge) — active deliverables
- **Updates** — change events (catalog/pricing/integration/etc.) for revenue-impact analysis
- **Shops** (MSO type only) — child shops

Cadences

A cadence is a recurring meeting schedule attached to an account. Define the frequency (weekly / biweekly / monthly / quarterly), day of week, and optional time. Folios:

- Auto-computes the next due date
- Sends in-app + browser-notification reminders 30m / 5m / start
- Auto-advances when a conversation is logged tied to it
- Surfaces the Cadence Hub when due

Cadences can also be **account-less** — a recurring 1:1 with your manager or a teammate, or an internal/leadership meeting (`cadence_scope='person'`). Action items from those meetings persist as **leadership tasks** — your own to-dos, filed under the 1:1 rather than forced onto a customer account — and show in that cadence's hub.

Cadence Hub

Pre-call command center for any account/cadence. Opens to:

- **Pip brief** — short pre-call read on the relationship state

- **Big Start Meeting button** — auto-creates a draft and opens full-screen meeting mode
- **Active Gauge projects** — expandable inline
- **Open items** — what's hanging
- **Follow-ups** — what's promised
- **History** — every prior conversation

Mobile uses a 4-tab segmented control (Notes / History / Tasks / Follow-ups) for compactness.

Full-screen meeting mode

Distraction-free conversation capture. Opens over the global chrome:

- Top bar with End & Summarize
- Collapsible sidebar with Pip brief, projects, open items, contacts
- Viewport-filling notepad with bullet-paste preservation
- Autosaves every 1.5s
- ESC to close

On End & Summarize, Pip generates a structured plan (action items, project updates, follow-ups) and the user reviews it in the [summarize-preview modal](#) before applying.

Discussed signal: during a meeting the user can tap any project card or open-item row in the sidebar to flag it as "discussed this meeting." Flagged cards show a teal ♦ Discussed chip. As the user types notes, the sidebar also auto-highlights any project or item whose title appears in the notes with a subtle glow — giving visual confirmation Pip will connect the dots before hitting summarize. Flagged IDs are passed to Pip at summarize time so Pip strongly prefers updating or closing those items instead of creating new duplicates. Rows in the plan modal that touch a flagged project or item show a ♦ Discussed badge so the user can see Pip acted on the signal.

Scheduling future one-off meetings

In addition to recurring cadences, you can pin a single upcoming meeting to a specific date/time. From the Calendar view:

- Click any empty day to open the schedule modal, or use the **+ Schedule Meeting** button in the header
- Pick the account (required), date, time, method, and an optional agenda note
- The meeting appears on the Calendar, Week, and List views with a distinct  chip (distinguished from cadence events)
- Reminders fire at 30m / 5m / start just like cadence meetings
- Today's scheduled meetings surface on the Home screen in a "Scheduled Today" card
- Opening the card on or after meeting day flips it into a live draft and opens the full-screen meeting mode so the conversation can be captured and summarized normally

Logging conversations (non-cadence)

For ad-hoc calls outside a cadence, the global "+ Conversation" button opens a modal: pick the account, pick the method, pick the date, then open the same full-screen meeting mode. The conversation files under the account's Meetings tab tagged **AD-HOC** .

Open items

Lightweight action-item tracking per account. Items can be:

- Manually added
- Auto-created by Pip from meeting notes (with **pip_created_at** marker so the V2 brain can detect post-creation edits)
- Carried forward across cadence meetings (open items show up in the next meeting's prep view)

Items support due dates, assignees (org members), and completion.

Quick Tasks

A floating tray for fast-capture tasks that don't belong to a specific account. 2-second log without leaving whatever screen you're on. Same underlying row shape as project

tasks.

Update Calendar

Manual change-event log per account. Categories: catalog, pricing, integration, product launch, training, promo, external event, other. Each entry has a date, title, description, owner, observed impact.

The value: when revenue dips on an account, you can correlate against the update timeline to see what changed and when. (Revenue surfaces themselves have been removed from the personal-mode build — see "Ripped" in CLAUDE.md — but the change-event log remains valuable.)

Inactive & Merge

No hard deletes for accounts or users. Everything is reversible:

- **Inactive flag** — soft archive. Greyed out in lists, excluded from Pip's "needs attention" counts. Can be reactivated.
- **Merge into...** — for acquisitions. Re-parents every child row (meetings, items, contacts, cadences, projects) from source → target, then marks source inactive. Linked back so history is preserved.

Pip — the AI field analyst

Pip is the differentiator. Embedded in every meeting, every account view, every brief.

Pip's personality

Loyal, slightly anxious field analyst. Same voice across all surfaces — meeting summaries, briefs, chat answers, error recovery messages. Personality is intentional: feels like a real teammate, not a generic chatbot.

Pip's modes

Mode	What it does	When it fires
summarize	Turns meeting notes into a structured plan (new items, updated items, closed items, project updates, follow-ups)	End & Summarize on a meeting
brief	Pre-call brief for an account (relationship state, what's at stake)	"Brief Me" button on Cadence Hub
ask	Conversational chat — ask anything about your accounts	Pip view (chat surface)
plan / extract	One-shot extraction of action items from short notes	Quick Touchpoint modal

Model tiering. Pip runs the right model for each task. The reasoning-heavy, user-facing, low-frequency surfaces — Ask Pip chat, meeting summarize, portfolio question generation, the daily brief, profile synthesis, and the QBR — run on **Sonnet 4.6**. High-volume or mechanical surfaces — per-account Brief Me, follow-up email drafting, terminology extraction, memory compression, and the leadership readout — run on **Haiku 4.5**. Each Sonnet surface has an environment override so its tier can be re-dialed without a code change. (See [ai-governance.md](#) for the full matrix.)

How Pip writes. Every Pip text surface — chat, the daily brief, meeting summaries, Brief Me, the cadence pre-call brief, and the QBR — renders structured output: a short headline, labeled sections, and scannable bullets with the important names in bold, instead of a wall of prose. Priority sections can carry one of Pip's own inline status glyphs (a small on-brand icon set — needs-now, keep-an-eye, good-news, cross-account-pattern, done) in place of generic emoji, so a brief reads at a glance. Email drafts stay plain prose (no markup) since they're meant to be sent as-is.

The V2 brain — Pip's learning loop

This is the differentiator that nothing else does. Pip doesn't just generate text; Pip **LEARNS** from how the user reacts to its output.

Four learning surfaces:

0. **User profile (folio_user_profile + profile_prose).** At first login, Pip interviews the user with 5 questions (role, company, portfolio, goals, communication style). Answers are synthesized into a 4–8 sentence `profile_prose` narrative injected into every Pip response — chat and summarize equally. Pip now knows who it's talking to, not just what accounts they have. Soft-gated: interview is skippable and resumable. Existing users see a dismissible HomeView card. Cost: ~\$0.002 once.

Phase 2 — drip questions. After onboarding, Pip keeps learning via a "Pip's Curious" card on the Home screen — one question at a time, inline textarea, never a modal. Soft cap of 5 per rolling 24h with no skip cooldown, so you can power through several in a sitting. Every question comes from real observed data — there is no generic "get-to-know-you" filler (that was removed; the well staying silent beats the well being fake). Sources:

- **Structural gaps (zero LLM cost):** contacts who appear in ≥ 3 meetings but have no role recorded, active accounts past 30 days with no objective, and empty profile slots post-onboarding.
- **Terminology lane (Lane C):** a daily Haiku scan of recent meeting notes surfaces proper nouns / brand names / codenames that appear ≥ 3 times but aren't a known account, contact, or glossary term. Each question is account-anchored ("you keep mentioning Fuse5 around John's Auto Parts — what is it?"). Answers write to `folio_pip_facts` so future briefs know the company's vocabulary automatically.
- **Portfolio generator (Lane D):** a low-frequency Sonnet pass reasons across the whole portfolio and writes a few genuinely insightful questions in Pip's voice. It self-skips (a DB count, no model call) whenever the queue already holds ≥ 5 , so it only spends when the queue has actually drained.

Answering ≥ 3 drip questions since the last synthesis triggers a background re-synthesis so `profile_prose` stays current. Settings → "Pip's Questions" has a global pause toggle and a completeness meter. Cost: gap detection = \$0; terminology + generator $\approx$ \$0.01–0.02/month; re-synthesis $\approx$ \$0.004 per batch.

1. **Correction log (pip_correction_log).** Every time the user declines a proposed action, edits the text of a Pip-created item, reassigns a task, or rejects a routing decision, the correction is captured. Pip reads the last 10 corrections

per account on every summarize call and is explicitly told "Chris has corrected these before — don't repeat."

2. **Lessons learned (`pip_account_state.lessons_learned`).** Every ~5 meetings, a cheap compression pass distills the recent correction log into a stable paragraph stored per account. Pip reads this paragraph instead of the raw log, keeping the read-back context small but the institutional memory permanent.
3. **Assignment hints (`pip_assignment_hints`).** When the user reassigns a task in the summarize-preview modal (e.g., changes assignee from "Mark" to "Sara" for invoice work), the pattern is captured. Pip routes the next similar task to Sara automatically. Hints scope: account-specific first, then org-wide ("Sara does invoice work everywhere").
4. **Glossary (`pip_glossary`).** Users can teach Pip terms, acronyms, and aliases (e.g., "ProParts was the legacy name for KSI Collision"). Pip respects glossary aliases when matching account names in notes.
5. **Structured suggestions ("Pip proposes, you approve").** When a drip question is *about* something structured — a contact's role, an account's objective, or a tool/system the account uses — the intent is attached to the question when it's created (Pip already knows which account or contact it's asking about). On answer, the card shows a pre-checked "Also save to..." toggle: keep it checked and the answer is written to the real field (`folio_contacts.title` , `folio_accounts.objective` , or the account's `systems` list); untick it to keep the answer as a plain fact. One tap, human in the loop — nothing mutates silently, and it never touches health or tier. The "systems" an account uses (e.g. "Fuse5 is their IMS") surface as chips on the account Overview and ride into every per-account Pip surface, so when "Fuse5" later shows up in raw notes Pip knows what it is instead of asking again. No per-answer model call — the structured write is free.

Cross-account routing. Pip can route a task from one account's meeting to a different account's plate. Example: "ACME meeting, but I need to follow up with KSI on the invoice feed" → task created on KSI, source backref points to the ACME meeting. The V2 brain handles the routing decision with a `target_account_id` field on each plan row. User can override in the preview modal.

Internal-meeting prior. When the meeting is on an internal-team account (`account_type = 'internal_team'`), Pip's default flips — it expects tasks to fan out to customer accounts, not stay on the internal account.

Pip — the summarize-preview modal

After Pip generates a meeting plan, the user reviews it before anything is applied. The modal shows every proposed action as a row with:

- Checkbox (default checked — uncheck to decline)
- Editable title (fix wording before applying)
- Account chip (override the routed account if Pip got it wrong)
- Assignee dropdown
- Due date picker
- "See source" link → expands the slice of the user's notes that triggered the row (editable; edited excerpts feed back as correction reasons)
- "Add missed item" button at the bottom for things Pip didn't catch

Grouped into Changes / New / Skipped. Yellow dot on low-confidence rows. Apply runs the selected rows through real DB writes. Cancel preserves the summarized meeting but applies nothing.

Pip — cost optimization

- **Model selection:** `claude-haiku-4-5-20251001` for default Pip operations. Sonnet only for high-synthesis paths (Brief Me).
- **Stacked prompt caching:** 4 cache breakpoints (system / glossary / roster / items+tasks). Multi-call sessions get ~70% cache hit rate.
- **Trivial-draft short-circuit:** drafts under 100 characters never call the API.
- **Per-user usage telemetry:** `folio_pip_usage` tracks every call (tokens, mode, timestamp); visible to user in Settings.
- **Rate limiting:** 20 req/min per user.

Pip — voice

- Speech-to-text via browser Web Speech API (free, browser-native).
- Text-to-speech via browser SpeechSynthesis today (free, but the voice quality is poor).
- Premium voice (OpenAI / Cartesia / ElevenLabs) on the wishlist.

Pip — portfolio intelligence ("chief of staff")

Pip sees your entire portfolio at once, not just one account at a time.

- **Daily account state snapshots** (`folio_account_snapshots`): computed once per day, client-side, from real signals — health status, days since last contact, open and overdue item counts, active and stuck Gauge project counts. Zero LLM cost; pure data derivation.
 - **Daily brief card on the Home screen**: one call per day, cached for the calendar day. Pip synthesises a morning read across the whole portfolio — what's due, what's at risk, recent wins — factoring in workload (overdue tasks, commitments, today's cadences), cross-account themes, tone trend, health momentum, and **off-cadence accounts vs. their own rhythm** (e.g. "you usually meet every ~3 weeks; it's been 45 days" — a personal baseline, not a generic cold threshold).
 - **Draft-ahead follow-ups**: Home surfaces meetings you summarised a couple days ago that still have no follow-up logged, with the follow-up email Pip already drafted ready to copy or send — no extra cost.
 - **Boss-ready leadership readout** and **pre-call standing agendas** (via the cadence brief) round out the proactive layer.
 - **Portfolio state utility** (`buildPortfolioState`): a compact text block summarising at-risk / watching accounts and stuck projects, injected into cross-portfolio Pip surfaces.
-

Gauge — projects & tasks

Gauge is the project-management module that lives under the Folios umbrella. Phase 1 is AM-facing; Phase 2 will expand to a full cross-functional PM tool.

Project types

Discrete projects. AM-owned, multi-step, has a finish line. The AM is the conductor; individual steps have assignees who may or may not be the AM. Example: an audit with 7 steps, 3 different people involved.

Standing projects. Admin-owned reactive queues. No sequence, no finish line. AMs (and Pip) drop tasks in; admin clears them. Example: an "Invoice Updates" inbox.

A project can be linked to multiple accounts — the account field is a searchable picker (search → pick → chip → search again), so a project that touches several accounts is filed under all of them.

Notes vs. status updates

Each project carries two distinct text surfaces:

- **Notes** — a durable scratchpad for context, decisions, and background that stays put.
- **Status updates** — an append-only, timestamped pulse log. Posting an update prepends a new entry (body + timestamp + author); the expanded card shows the latest pulse ("Updated 2h ago"), and the project edit screen shows the full history. Pip reads the latest update (plus the prior two) so briefs can say "All Star — latest: 'waiting on legal sign-off' (Jun 3)."

Custom fields

Each project defines its own custom field schema (text, longtext, number, date, dropdown, person, checkbox, URL). A "bones" preset provides defaults (Priority, Owner, Submission Date, Due Date, Description, Related Link) when starting blank.

Standing board view

Standing projects render as a kanban board inside the project row. One column per `task_status_columns` entry. Tasks drag between columns; admin processes left-to-right.

My Queue view

Flat task list across all projects the current user is assigned to. Live / Planning / All sub-filters. Optional group-by-project toggle.

Gauge — three lenses (queued, not yet built)

The next major Gauge build introduces three role-based default views set at invite time:

- **AM lens** — your accounts, your projects, your queue
- **Leader lens** — team rollup, who's stuck, what's overdue
- **Admin lens** — flat task queue, due dates, account chip per row

Pip's prompt branches per lens for role-appropriate framing. Same UI shape; different defaults and Pip context.

Full spec lives in CLAUDE.md → section 14 → Gauge V3.

Team & org layer

Multi-tenant ready from day 1:

- `folio_orgs` — org name, plan
- `folio_org_members` — user → org → role
- Roles: `owner | admin | member`
- All user-data tables carry `org_id` (or are linked via account)

- RLS policies scope to org membership where appropriate

Today: single-user pilot, no real org sharing in production. Architecture is ready when team mode lights up.

Light + dark themes

Full design system supporting both palettes. Tokens defined as CSS custom properties on `<html data-theme="...">`. Pre-mount inline script prevents flash-of-wrong-theme. Settings → Appearance for toggle.

Dark is the canonical palette for layout decisions; light is the spec'd translation. Light-mode-specific behaviors (hover lifts, mark pulses) are scoped via `[data-theme="light"]` so dark renders cleanly.

Mobile / PWA

Folios is mobile-first:

- Installable as a PWA on iOS and Android
 - Offline-capable shell (service worker)
 - Two redundant auto-update paths (controllerchange + version polling)
 - Safe-area-aware layouts
 - 16px minimum font on all inputs (prevents iOS auto-zoom)
 - Bottom nav: Home / Accounts / Calendar / Gauge
 - Settings/diagnostics/etc. tucked into the user menu
-

Observability

Self-monitoring built in:

- **folio_errors** — every React render error, network failure, Pip error, unhandled rejection lands here with stack + context. Visible in Settings → Diagnostics.
 - **Auto-recovery** — stale-chunk Lazy import failures (post-deploy edge case) detected and reloaded silently; marked auto-resolved.
 - **Connection status** — UI indicator when Supabase realtime drops.
 - **Render-time `timed()` helper** wraps hot paths for performance diagnostics.
-
-

What's NOT in Folios (deliberate)

- **Pipeline / revenue tracking** — ripped from the personal-mode build. Compliance blocks real revenue ingestion. Schema preserved for future re-build if corporate-data integration becomes possible. See CLAUDE.md → "Ripped (deliberate simplification)" for the inventory of what was pulled.
 - **Email integration** — no inbox sync, no auto-log from email threads. Pip operates on notes the user explicitly captures.
 - **CRM sync** — no Salesforce / HubSpot bidirectional sync today.
 - **Automatic transcription** — meeting notes are user-typed (or voice-dictated). No Otter-style auto-transcribe from a Zoom feed.
-
-

Roadmap snapshot

For the current shipping queue and wishlist, see CLAUDE.md → "Pending Updates" and "Feature Wishlist / Roadmap" sections (internal working doc). A presentation-facing

roadmap.md is planned.

Highlights of what's next:

- **Gauge V3** (three role-based lenses + unified `folio_tasks` table
 - project templates + Leader project view) — spec'd, queued.
 - **Premium TTS for Pip** (OpenAI / Cartesia / ElevenLabs) — wishlist.
 - **Pip's weekly clarifying questions** — spec'd, depends on more correction-log data first.
 - **Active session management UI** — for multi-device sign-out.
-
-

Contact

chris.vasconcellos97@gmail.com